Introduction to Advanced Data Analysis

Professor: Hammou El Barmi
Columbia University

Course Description: This course reviews and expands upon core topics in statistics through the study and practice of data analysis. It will be computationally intensive and we will use the R and or SAS languages and environments for statistical computing and graphics. I will assume that you had regression. The material that will be covered include: basic statistical procedures (including nonparametric procedures), linear regression, design of experiments, categorical data, survival analysis and multivariate analysis. Because of the nature of the course, I will use several books and they include

- The Statistical Sleuth: A Course in Methods of Data Analysis: Ramsey and Schafer, Duxbury
- Applied Linear Statistical Models: Kutner, Nachtsheim, Neter and Li (5th Edition), MaGraw Hill
- Methods of Multivariate Analysis, 2nd Edition, Alvin C. Rencher
- Categorical Data Analysis, 2nd Edition, Alain Agresti, Wiley
- The Statistical Analysis of Failure Time Data, 2nd Edition, J. D. Kalbfleisch and Ross L. Prentice, Wiley

Grading: Your final grade will be based on an in class exam, take home assignments (HWs) and a project. The homework will be assigned approximately every week and will collected a week later. Late assignments will not be accepted and they will result automatically in a zero. HWs submitted by email will NOT be graded and the lowest score on the homework assignments will be dropped. The exam will count for 40% of your final score and the HW will count for 30%. The project will count for the remaining 30%.

- Experiments are performed to generate data in order to make decisions.
- Much of the scientific knowledge about processes and systems is based on induction: reasoning from the specific to the general.
- Example(survey): Do NYC residents favor allowing guns in schools?
  - Specific cases: 100 NYC adults are called for a telephone survey
  - Inferential goal: get information about whether or not NYC residents favor allowing guns in schools.
- Example: Does hormone replacement improve health status in post-menopausal women?
  - Specific cases: the health status monitored in 1000 women over a 3-year period. Some took hormones, others did not.
  - Inferential goal: Determine if hormones improve the health of women not in the study.

- In general we are interested in how do the inputs of a process affect an output.
- The input variables can be divided into three categories:
  - controllable factors: measured and determined by scientist.
  - uncontrollable factors: measured but not determined by scientist.
  - noise factors: unmeasured, uncontrolled factors, often called experimental variability or error.

- For any interesting process, there are inputs such that:

$$\text{variability in input} \quad \Longrightarrow \quad \text{variability in output}$$

- If variability in x leads to variability y, we say x is a source of variation.
- Good design and analysis of experiments can identify sources of variation.

Information on how inputs affect the output can be gained from:

- Observational studies:
    - Input and output variables are observed from a pre-existing population.
    - It may be hard to say what is input and what is output.
- Controlled experiments: One or more input variables are controlled and manipulated by the experimenter to determine their effect on the output.

- Randomized, controlled, double-blind experiments
    - Randomization guards against selection bias.
    - Can eliminate correlation between x and y due to a different cause. aka a confounder. No causation without randomization
    - Ensure that the groups are comparable
    - Double-blind: minimizes bias in the response and in the evaluations of the experimental outcome
- Observational studies:
    - Assignment of experimental units to study groups is not done by the researcher
    - May lack advantages of controlled trials
    - May help establish association but not causation.
    - Can suggest good experiments to run, but can't definitively show causation.

When conducting a data analysis you need to ask the following questions:

- What is the objective of the analysis and/or the original experiment?
- What was the design of the study?
  - Randomized controlled or observational?
  - If a controlled trial, how were subjects assigned to the different groups?
  - Was the assignment process controlled by the investigator?
  - If an observational study: Are the groups comparable? What factors are confounded with treatment?
- What procedure would be appropriate for the data?
  - Exploratory data analysis techniques?
  - Inferential statistical techniques?
  - Model building?
- Implementation of analysis plan?
- Interpretation of Results?

- Preliminary look at data:
  - Evaluating data quality
  - Missing values
  - Outliers/Influential points
- Checking assumptions: Distributions, relationships, etc.
- Compute measures of location & dispersion

To do so, compute

- Numerical descriptive statistics (Measure of location, dispersion, skewness, kurtosis)
- Graphical descriptive statistics (histograms, box plots, stem and leaf plots, QQ plots)

- Identify research hypotheses to be tested.
- Choose a set of experimental units, which are the units to which treatments will be randomized.
- Choose a response/output variable.
- Determine potential sources of variation in response:
  - factors of interest
  - nuisance factors
- Decide which variables to measure and control:
  - treatment variables
  - potential large sources of variation in the units (blocking variables)
- Decide on the experimental procedure and how treatments are to be randomly assigned.