

COMS 4771 Machine Learning (Spring 2018)

Problem Set #3

Jianfu Yang - jy2863 Wenda Xu - wx2195 Fan Yang - fy2232
- wx2195@columbia.edu

March 27, 2018

Problem 1

(a) \mathbf{X} could be written as:

$$\begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1n} \\ x_{21} & x_{22} & \cdots & x_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ x_{d1} & x_{d2} & \cdots & x_{dn} \end{bmatrix}$$

\mathbf{y} could be written as:

$$[y_1 \ y_2 \ \cdots \ y_n]$$

Now transform \mathbf{X} to \mathbf{X}' :

$$\begin{bmatrix} x_{11} & x_{12} & \cdots & x_{1n} & \sqrt{\lambda\alpha} & 0 & \cdots & 0 \\ x_{21} & x_{22} & \cdots & x_{2n} & 0 & \sqrt{\lambda\alpha} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{d1} & x_{d2} & \cdots & x_{dn} & 0 & 0 & \cdots & \sqrt{\lambda\alpha} \end{bmatrix}$$

\mathbf{y} to \mathbf{y}' :

$$[y_1 \ y_2 \ \cdots \ y_n \ 0 \ 0 \ \cdots \ 0]$$

Then,

$$\|\mathbf{w}\mathbf{X}' - \mathbf{y}'\|_2^2 = \|\mathbf{w}\mathbf{X} - \mathbf{y}\|_2^2 + \lambda\alpha\|\mathbf{w}\|_2^2$$

We have,

$$\|\mathbf{w}\mathbf{X} - \mathbf{y}\|_2^2 + \lambda[\alpha\|\mathbf{w}\|_2^2 + (1 - \alpha)\|\mathbf{w}\|_1] = \|\mathbf{w}\mathbf{X}' - \mathbf{y}'\|_2^2 + \lambda(1 - \alpha)\|\mathbf{w}\|_1$$

Now, the new equivalent objective function is:

$$\arg \min_{\mathbf{w}} \|\mathbf{w}\mathbf{X}' - \mathbf{y}'\|_2^2 + \lambda(1 - \alpha)\|\mathbf{w}\|_1$$

It's a lasso regression with $\lambda' = \lambda(1 - \alpha)$.

(b) \mathbf{w} is a row vector, while \mathbf{x}_i is a column vector.

$$y_i = \mathbf{w}\mathbf{x}_i + \epsilon_i \sim \mathcal{N}(\mathbf{w}\mathbf{x}_i, \sigma^2)$$

$$w_j \sim \mathcal{N}(0, \tau^2)$$

$P(w_1, \dots, w_d | (x_1, y_1), \dots, (x_n, y_n))$ should consist of the likelihood and the prior. That is,

$$P = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \mathbf{w}\mathbf{x}_i)^2}{2\sigma^2}\right) \cdot \prod_{j=1}^d \frac{1}{\sqrt{2\pi\tau^2}} \exp\left(-\frac{w_j^2}{2\tau^2}\right)$$

$$\ln P = -\sum_{i=1}^n \frac{(y_i - \mathbf{w}\mathbf{x}_i)^2}{2\sigma^2} - \sum_{j=1}^d \frac{w_j^2}{2\tau^2} + \text{constant}$$

$$= -\frac{1}{2\sigma^2} \|\mathbf{w}\mathbf{X} - \mathbf{y}\|_2^2 - \frac{1}{2\tau^2} \|\mathbf{w}\|_2^2 + \text{constant}$$

$$\arg \max_{\mathbf{w}} P = \arg \max_{\mathbf{w}} \ln P = \arg \min_{\mathbf{w}} (-\ln P)$$

$$= \arg \min_{\mathbf{w}} \frac{1}{2\sigma^2} \|\mathbf{w}\mathbf{X} - \mathbf{y}\|_2^2 + \frac{1}{2\tau^2} \|\mathbf{w}\|_2^2$$

$$= \arg \min_{\mathbf{w}} \|\mathbf{w}\mathbf{X} - \mathbf{y}\|_2^2 + \frac{\sigma^2}{\tau^2} \|\mathbf{w}\|_2^2$$

$$= \arg \min_{\mathbf{w}} \|\mathbf{w}\mathbf{X} - \mathbf{y}\|_2^2 + \lambda \|\mathbf{w}\|_2^2$$

Thus, maximizing P is equivalent to minimizing the ridge optimization criterion. Proven.

Problem 2

(i)

$$\begin{aligned}
 D_{T+1}(i) &= \frac{D_{T+1}(i)}{\sum_i D_{T+1}(i)} \\
 &= \frac{1}{Z_T} D_T(i) e^{-\alpha_T y_i f_T(x_i)} \\
 &= \frac{1}{Z_T} e^{-\alpha_T y_i f_T(x_i)} \frac{1}{Z_{T-1}} e^{-\alpha_{T-1} y_i f_{T-1}(x_i)} D_{T-1}(i) \\
 &= \frac{1}{Z_T} e^{-\alpha_T y_i f_T(x_i)} \frac{1}{Z_{T-1}} e^{-\alpha_{T-1} y_i f_{T-1}(x_i)} \dots \frac{1}{Z_1} e^{-\alpha_1 y_i f_1(x_i)} D_1(i) \\
 &= \frac{1}{\prod_t Z_t} e^{-y_i \sum_t \alpha_t f_t(x_i)} D_1(i) \\
 &= \frac{1}{m} \frac{1}{\prod_t Z_t} e^{-y_i g(x_i)}
 \end{aligned}$$

Proven.

(ii) Firstly, since $D_{T+1}(i)$ is after normalizing,

$$\begin{aligned}
 \sum_i D_{T+1}(i) &= 1 \\
 \sum_i \frac{1}{m} \frac{1}{\prod_t Z_t} e^{-y_i g(x_i)} &= 1 \\
 \sum_i \frac{1}{m} e^{-y_i g(x_i)} &= \prod_t Z_t
 \end{aligned}$$

Then, using the fact that 0-1 loss is upper bounded by exponential loss:

$$\begin{aligned}
 err(g) &= \frac{1}{m} \sum_i \mathbf{1}[y_i \neq \text{sign}(g(x_i))] \\
 &\leq \frac{1}{m} \sum_i e^{-y_i g(x_i)} \\
 &= \prod_t Z_t \\
 err(g) &\leq \prod_t Z_t
 \end{aligned}$$

Proven.

(iii)

$$\begin{aligned}
Z_t &= \sum_i D_t(i) e^{-\alpha_t y_i f_t(x_i)} \\
&= \sum_i D_t(i) \mathbf{1}[y_i = f_t(x_i)] e^{-\alpha_t} + \sum_i D_t(i) \mathbf{1}[y_i \neq f_t(x_i)] e^{\alpha_t} \quad \left(\sum_i D_t(i) = 1 \right) \\
&= (1 - \sum_i D_t(i) \mathbf{1}[y_i \neq f_t(x_i)]) e^{-\alpha_t} + \sum_i D_t(i) \mathbf{1}[y_i \neq f_t(x_i)] e^{\alpha_t} \\
&= (1 - \epsilon_t) e^{-\frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}} + \epsilon_t e^{\frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}} \\
&= (1 - \epsilon_t) \sqrt{\frac{\epsilon_t}{1 - \epsilon_t}} + \epsilon_t \sqrt{\frac{1 - \epsilon_t}{\epsilon_t}} \\
&= 2\sqrt{\epsilon_t(1 - \epsilon_t)}
\end{aligned}$$

Proven.

(iv)

$$\begin{aligned}
err(g) &\leq \prod_t Z_t = \prod_t 2\sqrt{\epsilon_t(1 - \epsilon_t)} \quad (\epsilon_t = \frac{1}{2} - \gamma_t) \\
&= \prod_t 2\sqrt{(\frac{1}{2} - \gamma_t)(\frac{1}{2} + \gamma_t)} = \prod_t 2\sqrt{\frac{1}{4} - \gamma_t^2} \\
&= \prod_t \sqrt{1 - 4\gamma_t^2} \quad (1 + x \leq e^x \Rightarrow \sqrt{1 + x} \leq e^{\frac{1}{2}x}, \text{ for any } x \in \mathbb{R}) \\
&\leq \prod_t e^{-2\gamma_t^2} = e^{-2\sum_t \gamma_t^2}
\end{aligned}$$

Proven.

Problem 3

(i) Let \mathbf{b}' denotes the vector mapped from \mathbf{x}_i . Then:

$$b'_k = \sum_{j=1}^n A_{kj} x_{ij} \mod 2 \quad (1 \leq k \leq p)$$

Assume J contains the indexes of elements in \mathbf{x}_i which are 1. Since x_i is none-zero, $0 < |J| \leq n$. Then,

$$\sum_{j=1}^n A_{kj} x_{ij} = \sum_{j \in J} A_{kj}$$

Denote X as $\sum_{j \in J} A_{kj}$. Since A_{kj} obeys the Bernoulli distribution with $p = 0.5$, we have $X \sim B(J, \frac{1}{2})$. Then,

$$P(X) = C_J^X \left(\frac{1}{2}\right)^X \left(\frac{1}{2}\right)^{J-X} = C_J^X \left(\frac{1}{2}\right)^J$$

$b'_k = 0$ means X is even, $b'_k = 1$ means X is odd, that is,

$$P(b'_k = 0) = \left(\frac{1}{2}\right)^J \sum_{X \text{ is even}; X \leq |J|} C_J^X$$

$$P(b'_k = 1) = \left(\frac{1}{2}\right)^J \sum_{X \text{ is odd}; X \leq |J|} C_J^X$$

For binomial coefficient, we have

$$(1+x)^J = C_J^0 + C_J^1 x + \cdots + C_J^J x^J$$

$$(1+1)^J = C_J^0 + C_J^1 + \cdots + C_J^J = 2^J$$

$$(1-1)^J = C_J^0 - C_J^1 + C_J^2 - C_J^3 + \cdots = 0$$

Hence,

$$\sum_{X \text{ is even}; X \leq |J|} C_J^X = \sum_{X \text{ is odd}; X \leq |J|} C_J^X = 2^{J-1}$$

Now, we have,

$$P(b'_k = 0) = P(b'_k = 1) = \frac{1}{2}$$

Since \mathbf{b} has p elements, the probability of $\mathbf{b} = \mathbf{b}'$ is:

$$P = \left(\frac{1}{2}\right)^p$$

Proven.

- (ii) Denotes \mathbf{b} as the vector mapped from \mathbf{x}_i , and \mathbf{b}' as the vector mapped from \mathbf{x}_j . Then the probability of collision is the probability of $\mathbf{b} = \mathbf{b}'$:

$$P(\text{collision}) = \frac{1}{2^p}$$

(iii)

$$\begin{aligned}
 P(\text{no collisions}) &= \prod_{1 \leq i < j \leq n} P(\text{no collision between } \mathbf{x}_i \text{ and } \mathbf{x}_j) \\
 &= \left(1 - \frac{1}{2^p}\right)^{C_m^2} \\
 &= \left(1 - \frac{1}{2^p}\right)^{\frac{m^2-m}{2}} \quad (p = 2 \log_2 m) \\
 &= \left(1 - \frac{1}{m^2}\right)^{\frac{m^2-m}{2}} \quad (m \geq 1) \\
 &= f(m) \\
 \frac{d f(m)}{d m} &= \frac{\left(1 - \frac{1}{m^2}\right)^{\frac{m^2-m}{2}} ((2m^2 + m - 1) \log(1 - \frac{1}{m^2}) + 2)}{2(m+1)} \\
 &< 0 \quad (\text{for } m \geq 1)
 \end{aligned}$$

Hence, the minimum of $f(m)$ will be:

$$\min f(m) = \lim_{m \rightarrow \infty} f(m) = \lim_{m \rightarrow \infty} \left(1 - \frac{1}{m^2}\right)^{\frac{m^2-m}{2}} = \frac{1}{\sqrt{e}} \approx 0.6065$$

Thus, there will be no collision among \mathbf{x}_i with the probability of at least $\frac{1}{2}$.
Proven.

Problem 4

We firstly using linear mode (`sklearn.linear_model.LinearRegression`) to fit the data directly, and get an MAE around 102. Then we scale the data to have zero mean (`sklearn.preprocessing.StandardScaler`), and apply linear mode again, then get an MAE of 6.82.

To improve, we apply `sklearn.preprocessing.PolynomialFeatures` to the data, and again use linear mode (i.e., applying quadratic mode to the original data). This method decrease the MAE to 6.5.

We apply `sklearn.linear_model.Ridge`, no significant improvement observed.

Then we try to use `sklearn.svm.SVR` and `sklearn.kernel_ridge.KernelRidge` to fit the data. The former is too slow to get the result, while the latter uses too much memory.

We use `sklearn.linear_model.SGDRegressor`, setting the loss function as `epsilon_insensitive`, and get an MAE of 6.2.

We also try to use K-D tree and `sklearn.multioutput` to implement this regression problem as a classification problem, it doesn't make improvement.

Finally, we try to use MLP (`sklearn.neural_network.MLPRegressor`) to do the regression. We set the hidden layer size as `[25,5]`, using 'sgd' as solver, 'adaptive' as `step_size`, and finally get an MAE of around 6.00.