

# COMS 4771




## Perceptron and Kernelization

Nakul Verma

# Last time...


- Generative vs. Discriminative Classifiers
- Nearest Neighbor (NN) classification
- Optimality of  $k$ -NN
- Coping with drawbacks of  $k$ -NN
- Decision Trees
- The notion of overfitting in machine learning

# A Closer Look Classification



*Knowing the boundary is enough for classification*

# Linear Decision Boundary



*Assume binary classification  $y = \{-1, +1\}$   
(What happens in multi-class case?)*

# Learning Linear Decision Boundaries

$g$  = decision boundary

*d=1 case:*  $g(x) = w_1 x + w_0$

*general:*  $g(\vec{x}) = \vec{w} \cdot \vec{x} + w_0$

$f$  = linear classifier  $f(\vec{x}) := \begin{cases} +1 & \text{if } g(\vec{x}) \geq 0 \\ -1 & \text{if } g(\vec{x}) < 0 \end{cases}$


$$= \text{sign}(\vec{w} \cdot \vec{x} + w_0)$$

*# of parameters to learn in  $\mathbf{R}^d$ ?*


# Dealing with $w_0$

$$g(\vec{x}) = \vec{w} \cdot \vec{x} + w_0$$
$$= \underbrace{\begin{pmatrix} \vec{w} \\ w_0 \end{pmatrix}}_{\vec{w}'} \cdot \underbrace{\begin{pmatrix} \vec{x} \\ 1 \end{pmatrix}}_{\vec{x}'} \quad \text{bias}$$

$$g(\vec{x}') = \vec{w}' \cdot \vec{x}' \quad \text{homogeneous}$$




# The Linear Classifier



*A basic computational unit in a neuron*

# Can Be Combined to Make a Network



*Amazing fact:*

*Can approximate any  
smooth function!*

*An artificial neural network*



# How to Learn the Weights?

Given labeled training data (bias included):  $(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)$

Want:  $\vec{w}$ , which **minimizes** the training error, i.e.

$$\begin{aligned} \arg \min_{\vec{w}} \frac{1}{n} \sum_{i=1}^n \mathbf{1}[\text{sign}(\vec{w} \cdot \vec{x}_i) \neq y_i] \\ = \arg \min_{\vec{w}} \sum_{\substack{x_i \\ \text{s.t. } y_i = +1}} \mathbf{1}[\vec{x}_i \cdot \vec{w} < 0] + \sum_{\substack{x_i \\ \text{s.t. } y_i = -1}} \mathbf{1}[\vec{x}_i \cdot \vec{w} \geq 0] \end{aligned}$$

*How do we **minimize**?*

- Cannot use the standard technique (take derivative and examine the stationary points). Why?

*Unfortunately: NP-hard to solve or even approximate!*


# Finding Weights (Relaxed Assumptions)

Can we approximate the weights if we make reasonable assumptions?

*What if the training data is **linearly separable**?*

# Linear Separability

Say there is a **linear** decision boundary which can **perfectly separate** the training data



# Finding Weights

Given: labeled training data  $S = (\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)$

Want to determine: is there a  $\vec{w}$  which satisfies  $y_i(\vec{w} \cdot \vec{x}_i) \geq 0$  (for all  $i$ )

*i.e., is the training data linearly separable?*

Since there are  $d+1$  variables and  $|S|$  constraints, it is possible to solve efficiently it via a (constraint) optimization program. (How?)

*Can find it in a much **simpler** way!*

# The Perceptron Algorithm

Given: labelled training data  $S = (\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots (\vec{x}_n, y_n)$

Initialize  $\vec{w}^{(0)} = 0$


For  $t = 1, 2, 3, \dots$

If exists  $(\vec{x}, y) \in S$  s.t.  $\text{sign}(\vec{w}^{(t-1)} \cdot \vec{x}) \neq y$

$$\vec{w}^{(t)} \leftarrow \begin{cases} \vec{w}^{(t-1)} + \vec{x} & \text{if } y = +1 \\ \vec{w}^{(t-1)} - \vec{x} & \text{if } y = -1 \end{cases} = \vec{w}^{(t-1)} + y\vec{x}$$

(terminate when no such training sample exists)

# Perceptron Algorithm: Geometry




$$\text{sign}(\vec{w}^{(t-1)} \cdot \vec{x}) \neq +1$$

$$\vec{w}^{(t)} \leftarrow \vec{w}^{(t-1)} + \vec{x}$$

$$\text{sign}(\vec{w}^t \cdot \vec{x}) = +1$$

# Perceptron Algorithm: Geometry



$$\text{sign}(\vec{w}^{(t-1)} \cdot \vec{x}) \neq -1$$

$$\vec{w}^{(t)} \leftarrow \vec{w}^{(t-1)} - \vec{x}$$

$$\text{sign}(\vec{w}^t \cdot \vec{x}) = -1$$

# The Perceptron Algorithm

Input: labelled training data  $S = (\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)$

Initialize  $\vec{w}^{(0)} = 0$

For  $t = 1, 2, 3, \dots$

If exists  $(\vec{x}, y) \in S$  s.t.  $\text{sign}(\vec{w}^{(t-1)} \cdot \vec{x}) \neq y$

$$\vec{w}^{(t)} \leftarrow \begin{cases} \vec{w}^{(t-1)} + \vec{x} & \text{if } y = +1 \\ \vec{w}^{(t-1)} - \vec{x} & \text{if } y = -1 \end{cases} = \vec{w}^{(t-1)} + y\vec{x}$$

(terminate when no such training sample exists)

**Question:** Does the perceptron algorithm terminates? If so, when?



# Perceptron Algorithm: Guarantee

## Theorem (Perceptron mistake bound):

Assume there is a (unit length)  $\vec{w}^*$  that can separate the training sample  $S$  with margin  $\gamma$

Let  $R = \max_{\vec{x} \in S} \|\vec{x}\|$

Then, the perceptron algorithm will make at most  $T := \left(\frac{R}{\gamma}\right)^2$  mistakes.

*Thus, the algorithm will terminate in  $T$  rounds!*

*umm... but what about the generalization or the test error?*

# Proof

Key quantity to analyze:

How far is  $\vec{w}^{(t)}$  from  $\vec{w}^*$ ?

Suppose the perceptron algorithm makes a mistake in iteration  $t$ , then

$$\begin{aligned}\vec{w}^{(t)} \cdot \vec{w}^* &= (\vec{w}^{(t-1)} + y\vec{x}) \cdot \vec{w}^* \\ &\geq \vec{w}^{(t-1)} \cdot \vec{w}^* + \gamma\end{aligned}$$

$$\begin{aligned}\|\vec{w}^{(t)}\|^2 &= \|\vec{w}^{(t-1)} + y\vec{x}\|^2 \\ &= \|\vec{w}^{(t-1)}\|^2 + 2y(\vec{w}^{(t-1)} \cdot \vec{x}) + \|y\vec{x}\|^2 \\ &\leq \|\vec{w}^{(t-1)}\|^2 + R^2\end{aligned}$$

# Proof (contd.)

for all iterations  $t$

$$\vec{w}^{(t)} \cdot \vec{w}^* \geq \vec{w}^{(t-1)} \cdot \vec{w}^* + \gamma$$

$$\|\vec{w}^{(t)}\|^2 \leq \|\vec{w}^{(t-1)}\|^2 + R^2$$

So, after  $T$  rounds

$$T\gamma \leq \vec{w}^{(T)} \cdot \vec{w}^* \leq \|\vec{w}^{(T)}\| \|\vec{w}^*\| \leq R\sqrt{T}$$

Therefore: 
$$T \leq \left(\frac{R}{\gamma}\right)^2$$



# What Good is a Mistake Bound?

- It's an upper bound on the number of mistakes made by an **online algorithm** on an **arbitrary sequence** of examples  
*i.e. no i.i.d. assumption and not loading all the data at once!*
- Online algorithms with small mistake bounds can be used to develop classifiers with **good generalization error!**

# Other Simple Variants on the Perceptron

Voted perceptron

Average perceptron

Winnow



...

# Linear Classification

Linear classification simple,  
but... *when is real-data (even approximately) linearly separable?*


# What about non-linear decision boundaries?

Non linear decision boundaries are common:



# Generalizing Linear Classification

Suppose we have the following training data:



separable via a circular decision boundary

*d=2 case:*

$$g(\vec{x}) = w_1 x_1^2 + w_2 x_2^2 + w_0$$

*say, the decision boundary is some sort of ellipse*

*e.g. circle of radius r:*

$$w_1 = 1$$

$$w_2 = 1$$

$$w_0 = -r^2$$

*not linear in  $\vec{x}$  !*



# But $g$ is Linear in *some* Space!

$$g(\vec{x}) = w_1 x_1^2 + w_2 x_2^2 + w_0 \quad \text{non linear in } x_1 \text{ \& } x_2$$


$$= w_1 \chi_1 + w_2 \chi_2 + w_0 \quad \text{linear in } \chi_1 \text{ \& } \chi_2!$$

So if we apply a feature transformation on our data:

$$\phi(x_1, x_2) \mapsto (x_1^2, x_2^2)$$

Then  $g$  becomes linear in  $\phi$  - transformed feature space!

# Feature Transformation Geometrically



# Feature Transform for Quadratic Boundaries

**$\mathbf{R}^2$  case:** (*generic quadratic boundary*)

$$\begin{aligned} g(\vec{x}) &= w_1 x_1^2 + w_2 x_2^2 + w_3 x_1 x_2 + w_4 x_1 + w_5 x_2 + w_0 \\ &= \sum_{p+q \leq 2} w_{p,q} x_1^p x_2^q \end{aligned}$$

feature transformation:

$$\phi(x_1, x_2) \mapsto (x_1^2, x_2^2, x_1 x_2, x_1, x_2, 1)$$

**$\mathbf{R}^d$  case:** (*generic quadratic boundary*)

$$g(\vec{x}) = \sum_{i,j=1}^d \sum_{p+q \leq 2} w_{i,j}^{p,q} x_i^p x_j^q$$

*This captures all pairwise interactions between variables*

feature transformation:

$$\phi(x_1, x_2) \mapsto (x_1^2, x_2^2, \dots, x_d^2, x_1 x_2, \dots, x_{d-1} x_d, x_1, x_2, 1)$$

# Data is Linearly Separable in some Space!

## Theorem:

Given  $n$  labeled points  $(\vec{x}_1, y_1), (\vec{x}_2, y_2), \dots, (\vec{x}_n, y_n)$   $y_i = \{-1, +1\}$ ,  
there exists a feature transform where the data points are linearly separable.

*(this feature transform is sometimes called the Kernel transform)*

*the proof is almost trivial!*

# Proof

Given  $n$  points, consider the mapping into  $\mathbf{R}^n$ :

$$\phi(\vec{x}_i) \mapsto \begin{pmatrix} 0 \\ \vdots \\ 0 \\ y_i \\ 0 \\ \vdots \\ 0 \end{pmatrix} \quad \text{(zero in all coordinates except in coordinate } i \text{)}$$

Then, the decision boundary induced by linear weighting  $\vec{w}^* = \begin{pmatrix} 1 \\ \vdots \\ 1 \end{pmatrix}$  perfectly separates the input data!



# Transforming the Data into Kernel Space

Pros:

Any problem becomes **linearly separable**!

Cons:

What about **computation**? Generic kernel transform is typically  $\Omega(n)$

*Some useful kernel transforms map the input space into **infinite dimensional space**!*

What about **model complexity**?

*Generalization performance typically degrades with model complexity*

# The Kernel Trick (to Deal with Computation)

Explicitly working in generic Kernel space  $\phi(\vec{x}_i)$  takes time  $\Omega(n)$

But the **dot product** between two data points in kernel space can be computed relatively quickly

$$\phi(\vec{x}_i) \cdot \phi(\vec{x}_j) \quad \text{can compute fast}$$

**Example:** quadratic kernel transform for data in  $\mathbf{R}^d$

<i>explicit transform</i>	$O(d^2)$	$(x_1^2, \dots, x_d^2, x_1x_2, \dots, x_{d-1}x_d, x_1, x_2, 1)$
<i>dot products</i>	$O(d)$	$(1 + \vec{x}_i \cdot \vec{x}_j)^2$

RBF (radial basis function) kernel transform for data in  $\mathbf{R}^d$

<i>explicit transform</i>	infinite dimension!
<i>dot products</i>	$O(d)$

# The Kernel Trick

The trick is to perform classification in such a way that it **only accesses the data** in terms of **dot products** (so it can be done quicker)

**Example:** the 'kernel Perceptron'

*Recall:*  $\vec{w}^{(t)} \leftarrow \vec{w}^{(t-1)} + y\vec{x}$

*Equivalently*  $\vec{w} = \sum_{k=1}^n \alpha_k y_k \vec{x}_k$      $\alpha_k = \# \text{ of time mistake was made on } x_k$

*Thus, classification becomes*

$$f(\vec{x}) := \text{sign}(\vec{w} \cdot \vec{x}) = \text{sign}\left(\vec{x} \cdot \sum_{k=1}^n \alpha_k y_k \vec{x}_k\right) = \text{sign}\left(\sum_{k=1}^n \alpha_k y_k (\vec{x}_k \cdot \vec{x})\right)$$

*Only accessing data in  
terms of dot products!*



# The Kernel Trick: for Perceptron

*classification in original space:*

$$f(\vec{x}) = \text{sign}\left(\sum_{k=1}^n \alpha_k y_k (\vec{x}_k \cdot \vec{x})\right)$$

*If we were working in the transformed Kernel space, it would have been*

$$f(\phi(\vec{x})) = \text{sign}\left(\sum_{k=1}^n \alpha_k y_k (\phi(\vec{x}_k) \cdot \phi(\vec{x}))\right)$$

## Algorithm:

Initialize  $\vec{\alpha} = 0$

For  $t = 1, 2, 3, \dots, T$

If exists  $(\vec{x}_i, y_i) \in S$  s.t.  $\text{sign}\left(\sum_{k=1}^n \alpha_k y_k (\phi(\vec{x}_k) \cdot \phi(\vec{x}_i))\right) \neq y_i$

$\alpha_i \leftarrow \alpha_i + 1$

*implicitly working in  
non-linear kernel space!*

# The Kernel Trick: Significance

$$\sum_{k=1}^n \alpha_k y_k \left( \phi(\vec{x}_k) \cdot \phi(\vec{x}) \right)$$

*dot products are a measure of similarity*

***Can be replaced by any user-defined measure of similarity!***

*So, we can work in any user-defined non-linear space **implicitly** **without** the potentially heavy computational cost*

# What We Learned...

- Decision boundaries for classification
- Linear decision boundary (linear classification)
- The Perceptron algorithm
- Mistake bound for the perceptron
- Generalizing to non-linear boundaries (via Kernel space)
- Problems become linear in Kernel space
- The Kernel trick to speed up computation

# Questions?

# Next time...

Support Vector Machines (SVMs)!